

Question 3: RNN_ActionClassify

best performance on the Validation set: 87.5

1) baseline: 61.5

(project_layer): Linear(in_features=75, out_features=100, bias=True)


(recurrent_layer): LSTM(100, 100, batch_first=True)

(classify_layer): Linear(in_features=100, out_features=10, bias=True)

optimizer = torch.optim.SGD(model.parameters(), lr=1e-3, momentum=0.9)

criterion = nn.CrossEntropyLoss().type(dtype)

num_epochs = 75



2) 76.25

Increased trainable parameters in recurrent and classify layers,

Increased learning rate in SGD optimizer, added LR scheduler

(recurrent_layer): LSTM(75, 300, num_layers=3, batch_first=True)


(classify_layer): Linear(in_features=300, out_features=10, bias=True)

optimizer = torch.optim.SGD(model.parameters(), lr=1e-2, momentum=0.9)

scheduler = torch.optim.lr_scheduler.ReduceLROnPlateau(optimizer, patience=5, factor=0.1)

criterion = nn.CrossEntropyLoss().type(dtype)

num_epochs = 75




3) 83.75

Reduced trainable parameters in recurrent layer and added conv1d layer in classify layer. Rest same.

```
(recurrent_layer): LSTM(75, 200, batch_first=True)
(classify_layer): Sequential(
  (0): Conv1d(15, 10, kernel_size=(3,), stride=(1,))
  (1): BatchNorm1d(10, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
  (2): ReLU()
  (3): Flatten()
  (4): Linear(in_features=1980, out_features=10, bias=True)
)
```

```
optimizer = torch.optim.SGD(model.parameters(), lr=1e-2, momentum=0.9)
scheduler = torch.optim.lr_scheduler.ReduceLROnPlateau(optimizer, patience=5, factor=0.1)
criterion = nn.CrossEntropyLoss().type(dtype)
num_epochs = 75
```





4) 73.25

changed recurrent layer with dropout, bidirectional addition. Added one more lstm layer

```
(recurrent_layer): LSTM(75, 200, num_layers=2, batch_first=True, dropout=0.8,
bidirectional=True)
(r2): LSTM(75, 50, num_layers=2, batch_first=True, dropout=0.8)
(classify_layer): Sequential(
  (0): Linear(in_features=450, out_features=128, bias=True)
  (1): BatchNorm1d(128, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
  (2): LeakyReLU(negative_slope=0.01)
  (3): Linear(in_features=128, out_features=10, bias=True)
)
```

```
optimizer = torch.optim.SGD(model.parameters(), lr=1e-2, momentum=0.9)
scheduler = torch.optim.lr_scheduler.ReduceLROnPlateau(optimizer, patience=5, factor=0.1)
criterion = nn.CrossEntropyLoss().type(dtype)
num_epochs = 75
```





5) 83.75

From 3rd attempt, increased 1 layer in LSTM, used LeakyReLU in classify layer, used Adam optimizer

```
(recurrent_layer): LSTM(75, 200, num_layers=2, batch_first=True)
(classify_layer): Sequential(
  (0): Conv1d(15, 10, kernel_size=(3,), stride=(1,))
  (1): BatchNorm1d(10, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
  (2): LeakyReLU(negative_slope=0.01)
  (3): Flatten()
  (4): Linear(in_features=1980, out_features=10, bias=True)
)
```

```
optimizer = torch.optim.Adam(model.parameters(), lr=1e-3)
criterion = nn.CrossEntropyLoss().type(dtype)
num_epochs = 30
```





6) 85.5

From last attempt, again kept 1 layer in recurrent layer, used SGD optimizer (both same as 3rd attempt)

```
(recurrent_layer): LSTM(75, 200, batch_first=True)
(classify_layer): Sequential(
  (0): Conv1d(15, 10, kernel_size=(3,), stride=(1,))
  (1): BatchNorm1d(10, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
  (2): LeakyReLU(negative_slope=0.01)
  (3): Flatten()
  (4): Linear(in_features=1980, out_features=10, bias=True)
)
```

```
optimizer = torch.optim.SGD(model.parameters(), lr=1e-2, momentum=0.9)
scheduler = torch.optim.lr_scheduler.ReduceLROnPlateau(optimizer, patience=5, factor=0.1)
criterion = nn.CrossEntropyLoss().type(dtype)
num_epochs = 75
```



7) 87.5

Changed learning rate = 0.7×10^{-3} , added weight decay = 0.02

```
(recurrent_layer): LSTM(75, 200, batch_first=True)
(classify_layer): Sequential(
  (0): Conv1d(15, 10, kernel_size=(3,), stride=(1,))
  (1): BatchNorm1d(10, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
  (2): LeakyReLU(negative_slope=0.01)
  (3): Flatten()
  (4): Linear(in_features=1980, out_features=10, bias=True)
)
```

```
optimizer = torch.optim.SGD(model.parameters(), lr= $0.7 \times 10^{-3}$ , momentum=0.9, weight_decay=0.02)
```

```
scheduler = torch.optim.lr_scheduler.ReduceLROnPlateau(optimizer, patience=3, factor=0.1, verbose=True, min_lr=5e-5)
```

```
criterion = nn.CrossEntropyLoss().type(dtype)
```

```
num_epochs = 100
```

